

# HW 10

## KV – Store

Yue Wang, Hao Niu

### Test Parameters:

- Duration: 60 seconds per test
- Concurrency: 10 workers
- Total test time: ~16 minutes (all configurations)

## Performance Metrics Summary

### Read Latency (Mean)

Configuration	1%/99%	10%/90%	50%/50%	90%/10%
<b>W=5, R=1</b>	0.78ms	0.81ms	0.51ms	0.74ms
<b>W=1, R=5</b>	0.69ms	0.79ms	0.61ms	0.74ms
<b>W=3, R=3</b>	0.71ms	0.79ms	0.54ms	0.63ms
<b>Leaderless</b>	0.70ms	0.75ms	0.55ms	0.66ms

### Key Observations:

- All configurations show excellent read performance (< 1ms)
- Minimal variation between configurations
- Best performance: W=3, R=3 at 50%/50% (0.51ms)
- Worst performance: W=5, R=1 at 10%/90% (0.81ms)
- Difference is minimal (~0.3ms), all are very fast

### Write Latency (Mean)

<b>Configuration</b>	<b>1%/99%</b>	<b>10%/90%</b>	<b>50%/50%</b>	<b>90%/10%</b>
<b>W=5, R=1</b>	303.78ms	303.50ms	304.28ms	305.24ms
<b>W=1, R=5</b>	303.40ms	304.96ms	305.82ms	306.58ms
<b>W=3, R=3</b>	304.89ms	304.90ms	304.55ms	305.13ms
<b>Leaderless</b>	303.35ms	303.43ms	304.53ms	306.04ms

#### **Key Observations:**

- All configurations show consistent write latency (~303-307ms)
- Variation is minimal (~4ms range)
- Best performance: Leaderless at 1%/99% (303.35ms)
- Worst performance: W=1, R=5 at 90%/10% (306.58ms)
- Write latency dominated by replication delays, not strategy

## **Read Latency Percentiles (P95)**

<b>Configuration</b>	<b>1%/99%</b>	<b>10%/90%</b>	<b>50%/50%</b>	<b>90%/10%</b>
<b>W=5, R=1</b>	1.23ms	1.48ms	1.06ms	2.06ms
<b>W=1, R=5</b>	0.89ms	1.34ms	1.36ms	1.82ms
<b>W=3, R=3</b>	0.94ms	1.49ms	1.11ms	1.32ms
<b>Leaderless</b>	0.91ms	1.35ms	1.12ms	1.49ms

#### **Long Tail Analysis:**

- P95 values range from 0.89ms to 2.06ms
- Minimal long tail observed
- Highest P95: W=5, R=1 at 90%/10% (2.06ms)
- Lowest P95: W=1, R=5 at 1%/99% (0.89ms)

## **Write Latency Percentiles (P95)**

<b>Configuration</b>	<b>1%/99%</b>	<b>10%/90%</b>	<b>50%/50%</b>	<b>90%/10%</b>
<b>W=5, R=1</b>	305.71ms	305.43ms	307.22ms	308.24ms

<b>Configuration</b>	<b>1%/99%</b>	<b>10%/90%</b>	<b>50%/50%</b>	<b>90%/10%</b>
<b>W=1, R=5</b>	304.34ms	307.77ms	309.06ms	310.12ms
<b>W=3, R=3</b>	306.62ms	307.27ms	307.26ms	307.82ms
<b>Leaderless</b>	304.96ms	305.36ms	307.05ms	309.16ms

### Long Tail Analysis:

- P95 values range from 304.34ms to 310.12ms
- Very tight distribution (minimal tail)
- P95 is only ~2-5ms above mean
- No significant long tail for writes

## Stale Reads

**Observation:** 0 stale reads detected across all configurations.

### Analysis:

- All configurations: 0 stale reads
- This suggests replication completes before reads occur
- May require:
  - Higher concurrency (more simultaneous requests)
  - Longer test duration
  - More aggressive timing (reads immediately after writes)
  - Network delays or node failures

## Test Generator Explanation

### How It Works

Our load test generator uses a "**local-in-time**" key clustering algorithm:

1. **Key Organization:**
  - 1000 keys organized into 100 clusters (10 keys per cluster)
  - Each cluster represents related keys
2. **Active Cluster Tracking:**

- Maintains list of "active" clusters (accessed in last 5 seconds)
- Tracks last access time per cluster
- Expires clusters after 5 seconds of inactivity

**3. Selection Algorithm:**

- **80% probability:** Select from active clusters (recently accessed)
- **20% probability:** Start new cluster
- Within cluster: random key selection

**4. Temporal Locality:**

- Keys accessed recently stay "active"
- High probability (80%) of reusing active keys
- Creates bursts of activity on same keys

## How It Guarantees Local-In-Time Clustering

**1. 80/20 Probability Split:**

- 80% of requests target recently accessed keys (within 5 seconds)
- 20% explore new keys
- Ensures reads/writes to same key occur close together

**2. Active Cluster Window (5 seconds):**

- Keys accessed in last 5 seconds are "active"
- Multiple keys in same cluster accessed within this window
- Creates temporal clustering

**3. Cluster Expiration:**

- Prevents indefinite clustering
- Allows exploration while maintaining locality
- Realistic workload patterns

**4. Result:**

- Reads and writes to same key typically occur within 5-second windows
- Creates opportunities for stale reads
- Triggers "return most recent value" logic
- Demonstrates inconsistency windows

### Example:

Time 0.0s: Write key\_5 (cluster 0) → cluster 0 active

Time 0.1s: Read key\_7 (cluster 0) → 80% chance, cluster 0 active

Time 0.2s: Write key\_3 (cluster 0) → 80% chance, cluster 0 active

Time 0.3s: Read key\_9 (cluster 0) → 80% chance, cluster 0 active

Time 5.1s: Cluster 0 expires

Time 5.2s: Write key\_15 (cluster 1) → 20% chance, new cluster

## Configuration Performance Analysis

### Best Configuration by Workload Type

#### Read-Heavy (1% writes, 99% reads)

**Winner: W=1, R=5** (0.69ms read mean, 0.89ms P95)

#### Analysis:

- Fastest read latency
- Writes are rare, so write latency less important
- Eventual consistency acceptable for read-heavy workloads

**Alternative: Leaderless** (0.70ms read mean, 0.91ms P95)

- Nearly identical performance
- No single point of failure
- Good choice for distributed systems

#### Moderate Read (10% writes, 90% reads)

**Winner: Leaderless** (0.75ms read mean, 1.35ms P95)

#### Analysis:

- Best read performance
- Good write performance (303.43ms)
- Balanced approach

**Alternative: W=5, R=1** (0.81ms read mean, 1.48ms P95)

- Strong consistency
- Slightly slower reads

#### Balanced (50% writes, 50% reads)

**Winner: W=3, R=3** (0.54ms read mean, 1.11ms P95)

**Analysis:**

- Best read performance at this ratio
- Good write performance (304.55ms)
- Quorum provides fault tolerance
- Balanced consistency and performance

**Alternative: W=5, R=1** (0.51ms read mean, 1.06ms P95)

- Nearly identical performance
- Strong consistency

**Write-Heavy (90% writes, 10% reads)**

**Winner: W=3, R=3** (0.63ms read mean, 1.32ms P95)

**Analysis:**

- Best read performance
- Good write performance (305.13ms)
- Quorum balances consistency and performance

**Alternative: W=1, R=5** (0.74ms read mean, 1.82ms P95)

- Fast writes (but similar to others)
- Eventual consistency

## Key Findings and Discussion

### 1. Read Performance is Excellent Across All Configurations

**Finding:** All configurations show sub-millisecond read latency (0.51-0.81ms mean).

**Why:**

- Reads are local operations (R=1) or well-optimized (R=5, R=3)
- No significant coordination overhead
- Network latency is minimal (localhost)

**Implication:** Read performance is not a differentiating factor between configurations.

## 2. Write Latency is Consistent Across Strategies

**Finding:** All configurations show similar write latency (~303-307ms).

**Why:**

- Write latency dominated by replication delays (200ms + 100ms per node)
- Strategy differences are minimal compared to delay overhead
- Concurrent replication requests reduce perceived latency

**Implication:** Write strategy choice has minimal impact on latency in this setup.

## 3. Minimal Long Tail Observed

**Finding:** P95 and P99 values are close to mean, indicating minimal long tail.

**Why:**

- Consistent replication delays
- No network variability (localhost)
- No node failures or congestion

**Implication:** In production with network variability, long tails may be more pronounced.

## 4. No Stale Reads Detected

**Finding:** 0 stale reads across all configurations.

**Why:**

- Replication completes before reads occur
- Test timing may not catch inconsistency windows
- Client-side version tracking may need adjustment

**Implication:** Stale reads may require:

- Higher concurrency
- Longer test duration
- More aggressive timing
- Network delays

## 5. Time Intervals Confirm Local-In-Time Clustering

**Finding:** Time interval graphs show clustering of reads/writes to same key.

## **Why:**

- Test generator algorithm working as designed
- 80% probability of using active clusters
- 5-second active window creates temporal locality

**Implication:** Test generator successfully creates local-in-time patterns.

---

## **Recommendations**

### **For Read-Heavy Applications (1%/99%, 10%/90%)**

**Best Choice: W=1, R=5 or Leaderless**

- Fast read performance
- Eventual consistency acceptable
- No single point of failure (Leaderless)

### **For Balanced Applications (50%/50%)**

**Best Choice: W=3, R=3**

- Excellent read performance
- Good write performance
- Fault tolerance (quorum)
- Balanced consistency

### **For Write-Heavy Applications (90%/10%)**

**Best Choice: W=3, R=3 or W=1, R=5**

- Good performance for both operations
- W=1, R=5: Fast writes, eventual consistency
- W=3, R=3: Balanced approach with fault tolerance

## **General Recommendations**

### **1. Consistency Requirements:**

- Strong consistency needed → W=5, R=1
- Eventual consistency acceptable → W=1, R=5 or Leaderless

**2. Fault Tolerance:**

- Need fault tolerance → W=3, R=3 (quorum)
- No single point of failure → Leaderless

**3. Performance:**

- All configurations show excellent performance
- Differences are minimal
- Choose based on consistency and fault tolerance needs